



A Deep Dive into Monte Carlo Tree Search


Brian Lee (brianklee@google.com)



www.moderndescartes.com

Pycon 2018

About Me


- Software engineer at Verily Life Sciences (an Alphabet company).
- Local organizer for Go events.
- First programming project was a Django website to manage Go tournaments (2012).
- More recently: Go AI (2016)






MuGo
(github.com/brilee/MuGo)

¹David Silver, et. al. Mastering the game of Go with deep neural networks and tree search. *Nature*, 529(7587):484–489, January 2016



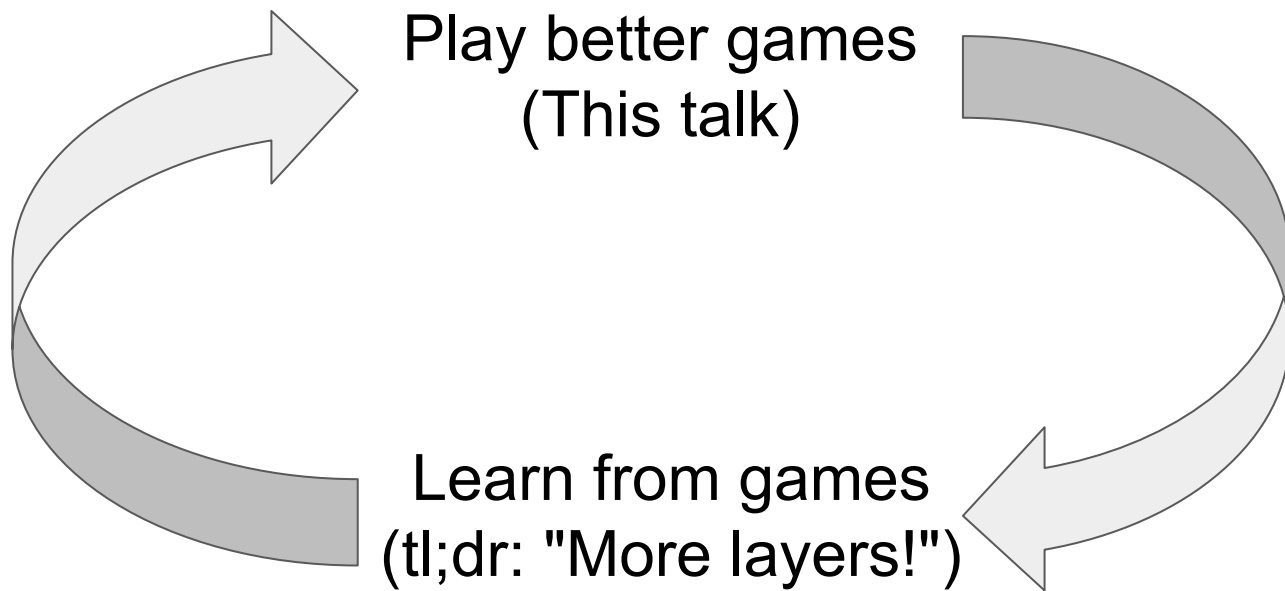
MuGo
(github.com/brilee/MuGo)

¹David Silver, et. al. Mastering the game of go without human knowledge. *Nature*, 550:354– 359, 2017.



github.com/tensorflow/minigo


AlphaGoZero, summarized



Come see MiniGo in action!



A simpler problem

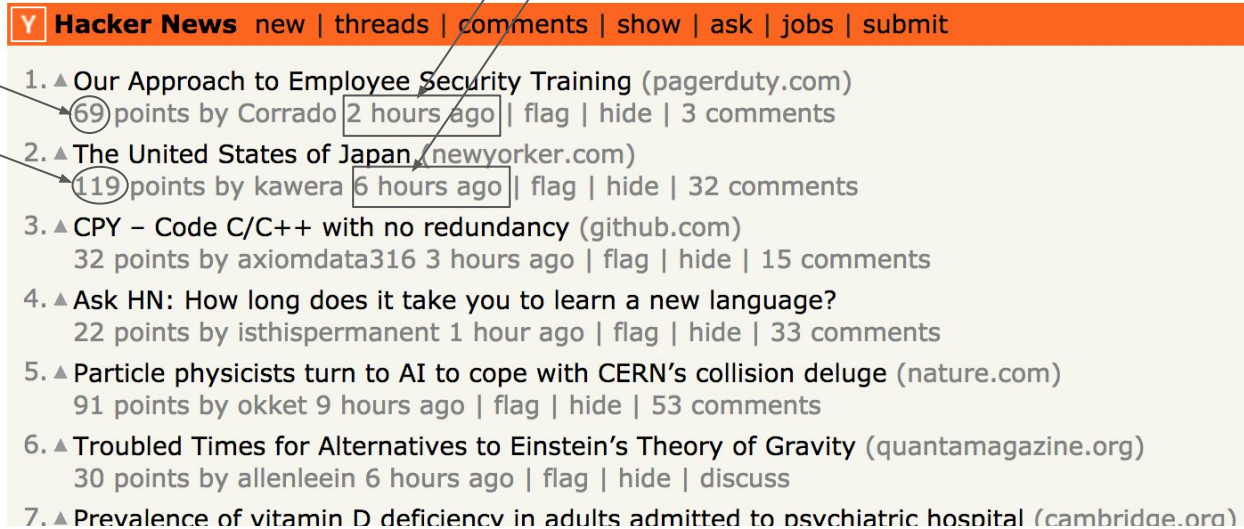
 **Hacker News** [new](#) | [threads](#) | [comments](#) | [show](#) | [ask](#) | [jobs](#) | [submit](#)

1. ▲ Our Approach to Employee Security Training ([pagerduty.com](#))
69 points by Corrado 2 hours ago | [flag](#) | [hide](#) | 3 comments
2. ▲ The United States of Japan ([newyorker.com](#))
119 points by kawera 6 hours ago | [flag](#) | [hide](#) | 32 comments
3. ▲ CPY – Code C/C++ with no redundancy ([github.com](#))
32 points by axiomdata316 3 hours ago | [flag](#) | [hide](#) | 15 comments
4. ▲ Ask HN: How long does it take you to learn a new language?
22 points by isthispermanent 1 hour ago | [flag](#) | [hide](#) | 33 comments
5. ▲ Particle physicists turn to AI to cope with CERN's collision deluge ([nature.com](#))
91 points by okket 9 hours ago | [flag](#) | [hide](#) | 53 comments
6. ▲ Troubled Times for Alternatives to Einstein's Theory of Gravity ([quantamagazine.org](#))
30 points by allenleein 6 hours ago | [flag](#) | [hide](#) | [discuss](#)
7. ▲ Prevalence of vitamin D deficiency in adults admitted to psychiatric hospital ([cambridge.org](#))

Solution:
Upper Confidence Bounds

Uncertainty

Best guess
at quality



The image shows a screenshot of the Hacker News website. At the top is an orange navigation bar with the text 'Y Hacker News' followed by links: 'new | threads | comments | show | ask | jobs | submit'. Below this is a list of seven news items. Annotations are present: an arrow from the text 'Best guess at quality' points to the number of points for the first item (69); another arrow from the same text points to the number of points for the second item (119); and an arrow from the word 'Uncertainty' points to the time '2 hours ago' for the first item and '6 hours ago' for the second item.

Y Hacker News new | threads | comments | show | ask | jobs | submit


1. ▲ Our Approach to Employee Security Training (pagerduty.com)
69 points by Corrado 2 hours ago | flag | hide | 3 comments
2. ▲ The United States of Japan (newyorker.com)
119 points by kawera 6 hours ago | flag | hide | 32 comments
3. ▲ CPY – Code C/C++ with no redundancy (github.com)
32 points by axiomdata316 3 hours ago | flag | hide | 15 comments
4. ▲ Ask HN: How long does it take you to learn a new language?
22 points by isthispermanent 1 hour ago | flag | hide | 33 comments
5. ▲ Particle physicists turn to AI to cope with CERN's collision deluge (nature.com)
91 points by okket 9 hours ago | flag | hide | 53 comments
6. ▲ Troubled Times for Alternatives to Einstein's Theory of Gravity (quantamagazine.org)
30 points by allenleein 6 hours ago | flag | hide | discuss
7. ▲ Prevalence of vitamin D deficiency in adults admitted to psychiatric hospital (cambridge.org)

Ranking = quality + upper confidence
bound

To learn more:
multi armed bandit


UCT = Upper Confidence
bounds (applied to Trees)

Fig. 2 from AlphaGoZero paper




David Silver, et. al. Mastering the game of go without human knowledge. *Nature*, 550:354– 359, 2017.

Step 1: Select



Q: Our best guess at the move's value
U: The upper confidence bound


Step 2: Evaluate, Expand



P: Neural net's move probabilities (Array of floats)


V: Neural net's position evaluation (Float)

Step 3: Backup



$Q = \text{Average of all children's evaluation } V$

Step 4: repeat



David Silver, et. al. Mastering the game of go without human knowledge. *Nature*, 550:354– 359, 2017.

Implementing UCT in Python

All code is online at
github.com/brilee/python_uct

The UCT Algorithm

```
def UCT_search(game_state, num_reads):  
    root = UCTNode(game_state)  
    for _ in range(num_reads):  
        leaf = root.select_leaf()  
        child_priors, value_estimate = NeuralNet.evaluate(leaf.game_state)  
        leaf.expand(child_priors)  
        leaf.backup(value_estimate)  
    return max(root.children.items(), key=lambda item: item[1].number_visits)
```


class UCTNode

```
class UCTNode():  
    def __init__(self, game_state, parent=None, prior=0):  
        self.game_state = game_state  
        self.is_expanded = False  
        self.parent = parent # Optional[UCTNode]  
        self.children = {} # Dict[move, UCTNode]  
        self.prior = prior # float  
        self.total_value = 0 # float  
        self.number_visits = 0 # int
```

Step 1: Select

```
def Q(self): # returns float
    return self.total_value / (1 + self.number_visits)
```

```
def U(self): # returns float
    return (math.sqrt(self.parent.number_visits) * self.prior /
            (1 + self.number_visits))
```




```
def UCT_search(game_state, num_reads):
    root = UCTNode(game_state)
    for _ in range(num_reads):
        leaf = root.select_leaf()
        child_priors, value_estimate = NeuralNet.evaluate(leaf.game_state)
        leaf.expand(child_priors)
        leaf.backup(value_estimate)
    return max(root.children.items(), key=lambda item: item[1].number_visits)
```

Step 1: Select

```
def best_child(self):  
    return max(self.children.values(),  
               key=lambda node: node.Q() + node.U())
```

```
def select_leaf(self):  
    current = self  
    while current.is_expanded:  
        current = current.best_child()  
    return current
```




```
def UCT_search(game_state, num_reads):  
    root = UCTNode(game_state)  
    for _ in range(num_reads):  
        leaf = root.select_leaf()  
        child_priors, value_estimate = NeuralNet.evaluate(leaf.game_state)  
        leaf.expand(child_priors)  
        leaf.backup(value_estimate)  
    return max(root.children.items(), key=lambda item: item[1].number_visi
```

Step 2: Evaluate, expand

```
def expand(self, child_priors):
    self.is_expanded = True
    for move, prior in enumerate(child_priors):
        self.add_child(move, prior)
```


```
def add_child(self, move, prior):
    self.children[move] = UCTNode(
        self.game_state.play(move), parent=self, prior=prior)
```



```
def UCT_search(game_state, num_reads):
    root = UCTNode(game_state)
    for _ in range(num_reads):
        leaf = root.select_leaf()
        child_priors, value_estimate = NeuralNet.evaluate(leaf.game_state)
        leaf.expand(child_priors)
        leaf.backup(value_estimate)
    return max(root.children.items(), key=lambda item: item[1].number_visits)
```


Step 3: Backup

```
def backup(self, value_estimate: float):  
    current = self  
    while current.parent is not None:  
        current.number_visits += 1  
        current.total_value += (value_estimate  
                                * self.game_state.to_play)  
        current = current.parent
```



```
def UCT_search(game_state, num_reads):  
    root = UCTNode(game_state)  
    for _ in range(num_reads):  
        leaf = root.select_leaf()  
        child_priors, value_estimate = NeuralNet.evaluate(leaf.game_state)  
        leaf.expand(child_priors)  
        leaf.backup(value_estimate)  
    return max(root.children.items(), key=lambda item: item[1].number_visits)
```

Performance of simple code

For num_reads = 10000...

Memory usage: 1.8GB

Time: 30 sec

Numpy: SIMD In Python

```
>>> nodes = [(0.7, 0.1), (0.3, 0.3),  
(0.4, 0.2)]  
>>> q_plus_u = [_1 + _2 for _1, _2 in  
nodes]  
>>> q_plus_u  
[0.8, 0.6, 0.6]  
>>> max(range(len(q_plus_u)),  
key=lambda i: q_plus_u[i])  
0
```

```
>>> import numpy as np  
>>> q = np.array([0.7, 0.3, 0.4])  
>>> u = np.array([0.1, 0.3, 0.2])  
>>> q_plus_u = q + u  
>>> q_plus_u  
array([0.8, 0.6, 0.6])  
>>> np.argmax(q_plus_u)  
0
```

```
class UCTNode():
    def __init__(self, game_state,
                  parent=None, prior=0):
        self.game_state = game_state
        self.is_expanded = False
        self.parent = parent # Optional[UCTNode]
        self.children = {} # Dict[move, UCTNode]
        self.prior = prior # float
        self.total_value = 0 # float
        self.number_visits = 0 # int
```

```
class UCTNode():
    def __init__(self, game_state,
                  move, parent=None):
        self.game_state = game_state
        self.move = move
        self.is_expanded = False
        self.parent = parent # Optional[UCTNode]
        self.children = {} # Dict[move, UCTNode]
        self.child_priors = np.zeros(
            [362], dtype=np.float32)
        self.child_total_value = np.zeros(
            [362], dtype=np.float32)
        self.child_number_visits = np.zeros(
            [362], dtype=np.float32)
```

```
@property
def number_visits(self):
    return self.parent.child_number_visits[self.move]

@number_visits.setter
def number_visits(self, value):
    self.parent.child_number_visits[self.move] = value

@property
def total_value(self):
    return self.parent.child_total_value[self.move]

@total_value.setter
def total_value(self, value):
    self.parent.child_total_value[self.move] = value
```

```
def Q(self): # returns float
    return self.total_value /
        (1 + self.number_visits)
```

```
def U(self): # returns float
    return (math.sqrt(self.parent.number_visits)
            * self.prior / (1 + self.number_visits))
```

```
def best_child(self):
    return max(self.children.values(),
               key=lambda node: node.Q() + node.U())
```

```
def child_Q(self): # returns np.array
    return self.child_total_value /
        (1 + self.child_number_visits)
```

```
def child_U(self): # returns np.array
    return math.sqrt(self.number_visits) *
        (self.child_priors / (1 + self.child_number_visits))
```

```
def best_child(self):
    return np.argmax(self.child_Q() + self.child_U())
```

Performance of optimized code

For num_reads = 10000...

Memory usage: 90MB (20x improvement)

Time: 0.8 sec (40x improvement)

Performance of optimized code

For time = 30 sec

Memory usage: 1.7GB (Same as before)

Num_reads: 300,000 (30x improvement)

Other components of a battle-tested
UCT implementation...

- Handle illegal moves
- When game end condition is triggered, use actual scoring, not network
- Impose a move limit to prevent long games
- Subtree reuse
- Pondering
- Parent-Q initialization
- Tuning relative weighting of Q, U
- Virtual Losses

What about
the GPU?

The UCT Algorithm

```
def UCT_search(game_state, num_reads):  
    root = UCTNode(game_state)  
    for _ in range(num_reads):  
        leaf = root.select_leaf()  
        child_priors, value_estimate = NeuralNet.evaluate(leaf.game_state)  
        leaf.expand(child_priors)  
        leaf.backup(value_estimate)  
    return max(root.children.items(), key=lambda item: item[1].number_visits)
```

Virtual losses to the rescue!

```
@@ -27,6 +27,8 @@ def best_child(self):
```

```
    def select_leaf(self):
```

```
        current = self
```

```
        while current.children:
```

```
+            current.number_visits += 1
```

```
+            current.total_value -= 1
```

```
            current = current.best_child()
```

```
        return current
```

```
@@ -38,8 +40,7 @@ def expand(self, child_priors: List[fl
```

```
    def backup(self, value_estimate: float):
```

```
        current = self
```

```
        while current.parent is not None:
```

```
-            current.number_visits += 1
```

```
-            current.total_value += value_estimate
```

```
+            current.total_value += value_estimate + 1
```

```
            current = current.parent
```

Virtual losses:

Up to 50x faster with right hardware

Brian Lee

(brianklee@google.com)

www.moderndescartes.com